

ALFIDO TECH

TASK 2===

CODE=

Import required libraries

```
import pandas as pd

from sklearn.model_selection

import train_test_split

from sklearn.preprocessing

import LabelEncoder, StandardScaler

from sklearn.linear_model

import LogisticRegression

from sklearn.tree

import DecisionTreeClassifier

from sklearn.metrics

import accuracy_score, confusion_matrix, classification_report
```

=== Step 1: Create sample data ===

```
data = { 'Name': ['Alice', 'Bob', 'Carol', 'David', 'Eva'],
         'Age': [25, 30, None, 35, 28],
         'Salary': [50000, None, 60000, 70000, 55000],
         'Gender': ['Female', 'Male', 'Female', 'Male', 'Female'],
         'Purchased': ['Yes', 'No', 'Yes', 'No', 'Yes'] }
```

```
df = pd.DataFrame(data)
```

=== Step 2: Preprocess ===

```
df['Age'].fillna(df['Age'].mean(), inplace=True)
```

```
df['Salary'].fillna(df['Salary'].mean(), inplace=True)
```

Encode categorical columns

```
label_enc = LabelEncoder()
```

```
for col in ['Name', 'Gender', 'Purchased']:
```

```
    df[col] = label_enc.fit_transform(df[col])
```

=== Step 3: Split into features (X) and target (y) ===

```
X = df.drop(columns=['Purchased'])
```

```
y = df['Purchased'] # keep as integer labels
```

=== Step 4: Scale features only (not the target!) ===

```
scaler = StandardScaler()
```

```
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
```

=== Step 5: Train-test split ===

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,  
random_state=42)
```

=== Step 6: Logistic Regression Model ===

```
log_model = LogisticRegression()  
  
log_model.fit(X_train, y_train)  
  
log_pred = log_model.predict(X_test)  
  
print("📊 Logistic Regression Results:")  
  
print("Accuracy:", accuracy_score(y_test, log_pred))  
  
print("Confusion Matrix:\n", confusion_matrix(y_test, log_pred))  
  
print("Classification Report:\n", classification_report(y_test, log_pred))
```

=== Step 7: Decision Tree Model ===

```
tree_model = DecisionTreeClassifier()  
  
tree_model.fit(X_train, y_train)  
  
tree_pred = tree_model.predict(X_test)  
  
print("\n🌳 Decision Tree Results:")  
  
print("Accuracy:", accuracy_score(y_test, tree_pred))  
  
print("Confusion Matrix:\n", confusion_matrix(y_test, tree_pred))  
  
print("Classification Report:\n", classification_report(y_test, tree_pred))
```

OUTPUT=====

```
File Edit Selection View Go Run Terminal Help
task_2.ipynb alfidotask1.ipynb Python Debugger: Core...

C:\Users\mithy> nithy> alfidotask1.ipynb> # Import required libraries
Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ...
base (Python 3.12.7)

# === Step 7: Decision Tree Model ===
tree_model = DecisionTreeClassifier()
tree_model.fit(X_train, y_train)
tree_pred = tree_model.predict(X_test)

print("\n Decision Tree Results:")
print("Accuracy:", accuracy_score(y_test, tree_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, tree_pred))
print("Classification Report:\n", classification_report(y_test, tree_pred))

C:\Users\mithy\AppData\Local\Temp\ipykernel_11856\3207397328.py:21: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead.

df['Age'].fillna(df['Age'].mean(), inplace=True)
C:\Users\mithy\AppData\Local\Temp\ipykernel_11856\3207397328.py:22: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead.

df['Salary'].fillna(df['Salary'].mean(), inplace=True)
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
```

```
File Edit Selection View Go Run Terminal Help
task_2.ipynb alfidotask1.ipynb Python Debugger: Core...

C:\Users\mithy> nithy> alfidotask1.ipynb> # Import required libraries
Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline ...
base (Python 3.12.7)

Logistic Regression Results:
Accuracy: 0.0
Confusion Matrix:
[[0 1]
 [0 0]]
Classification Report:
      precision    recall  f1-score   support

     0       0.00      0.00      0.00         1.0
     1       0.00      0.00      0.00         0.0

 accuracy          0.00      0.00      0.00         1.0
 macro avg          0.00      0.00      0.00         1.0
 weighted avg        0.00      0.00      0.00         1.0

Decision Tree Results:
Accuracy: 0.0
Confusion Matrix:
[[0 1]
 [0 0]]
Classification Report:
      precision    recall  f1-score   support

     0       0.00      0.00      0.00         1.0
     1       0.00      0.00      0.00         0.0

 accuracy          0.00      0.00      0.00         1.0
 macro avg          0.00      0.00      0.00         1.0
 weighted avg        0.00      0.00      0.00         1.0

Output is truncated. View as a scrollable element or open in a text editor. Adjust call output settings.
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
C:\Users\mithy\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
warn_prf(average, modifier, f"(metric.capitalize()) is", len(result))
```